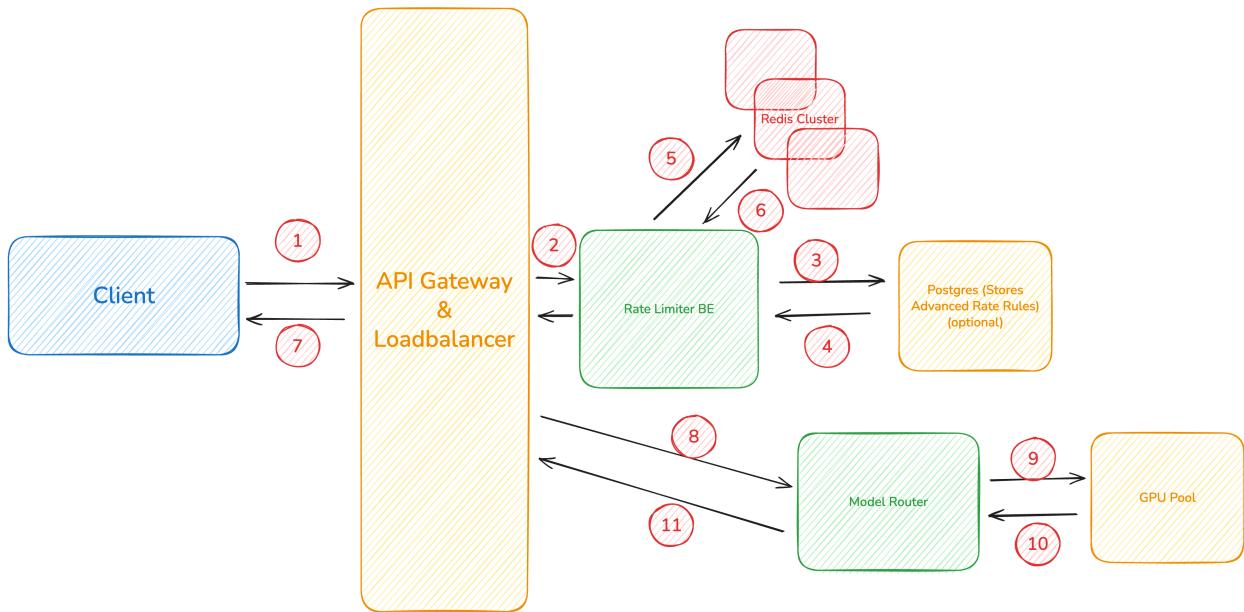


### 2.1 Architecture Overview

Client → API Gateway → Rate Limiter → Model Router → GPU Pool



### Purpose of rate limiter in AI serving

- Stops abusive clients from overloading GPUs.
- Enforces tenant-level SLAs.
- Prioritizes internal / premium traffic (QoS tiers).

### 2.2 Components

#### Frontend (React)

- Sends requests to backend
- Displays decision

#### Rate Limiter Backend (FastAPI)

- Implements /rate-limit/check
- Uses Redis Sliding Window

#### Redis Cluster

- Stores sliding window ZSETs
- May be sharded across multiple nodes

### Postgres

- Stores advanced rate policies:
  - per tenant
  - per API key
  - per model tier

### 2.3 Data Flow

1. UI/API sends rate-limit check
2. Backend builds Redis key
3. Lua script prunes old entries + counts + inserts atomically
4. Backend returns allowed/block
5. Client decides whether to call AI inference API
6. If allowed → request hits GPU pool

### 2.4 Key Design Choice: Sliding Window Log

Compared to fixed window or token bucket, Sliding Window Log:

- Gives **fairer** enforcement
- Smooths bursts
- Ideal for AI inference where each call may be expensive

### 2.5 Distributed Design

- Many FastAPI instances → one Redis cluster
- Lua script ensures atomic ops
- Keys hashed across Redis shards
- No cross-shard transactions required

### 2.6 AI-Specific Extensions

#### Protecting GPU pools

- Add **model-global** rate limits:
- rl:model:{modelId}:global
- Prevents GPU saturation beyond capacity.

#### Token-based billing

- Track cost or tokens:
- rl:tokens:{userId}:{modelId}
- Block if exceeding token/hour or cost/day budgets.

### QoS tiers

- Tier-based limits:
  - Internal: very high
  - Enterprise: medium
  - Free: strict
- Keys:
  - rl:qos:{tier}:{modelId}

### 2.6 For very high daily active user count and throughput

#### 1. We might need to shard Redis due to two things:

1. **Memory** – how many timestamps we store
2. **Throughput** – how many ops/sec we need

Because we're using a **Sliding Window Log** with:

- limit = 100 requests/hour (per key)
- each request = 1 ZSET entry
- old entries are purged + TTL removes idle keys  
→ **each active key is bounded to ~100 entries**

#### Rough memory math

Approximate sizing:

- Assume ~100 bytes per ZSET entry (timestamp, metadata, Redis overhead).
- For 100 requests/hour per (user, model):

Per key:

- $100 \text{ entries} \times 100 \text{ bytes} \approx 10 \text{ KB}$

Now say we have 1,000,000 active (user, model) pairs:

- $1,000,000 \text{ keys} \times 10 \text{ KB} \approx 10,000,000 \text{ KB} \approx 10 \text{ GB}$

So:

- **100k active keys** → ~1 GB
- **1M active keys** → ~10 GB

On a typical 8–16 GB Redis instance, **1M+ active keys** starts to be tight (plus overhead for other data), and sharding / clustering becomes attractive.

### 2. Throughput considerations

Each rate-limit check does:

- ZREMRANGEBYSCORE
- ZCARD
- ZADD
- EXPIRE

Wrapped in a WATCH/MULTI/EXEC pipeline.

A single good Redis node can often handle **100k+ ops/sec**, but that depends on instance size, network, and whether other workloads are sharing it.

Rule of thumb:

- **A few thousand QPS** → single Redis node is totally fine.
- **Tens of thousands QPS sustained**, plus other workloads → **we need to start planning for Redis Cluster or managed sharding.**

### 3. The current design already supports sharding

**The current key design already supports sharding later:**

```
rl:user:{userId}:model:{modelId}  
rl:tenant:{tenantId}:model:{modelId}  
rl:apikey:{apiKey}:model:{modelId}  
...
```

Redis Cluster will automatically hash these keys across slots/nodes; you do **not** need to change the key format later.

### 4. What to do when we *do* need sharding

When we grow, we have a few options:

#### Option A – Redis Cluster (or AWS ElastiCache / Azure Cache)

- Enable Redis Cluster or use managed service with cluster mode.
- Point our client (redis-py) to the cluster endpoints.
- Because each rate-limit check touches **one key**, we don't need cross-slot transactions; each request lives on one shard.

Your current logic still works:

```
key = f"rl:user:{userId}:model:{modelId}"  
rate_limiter.check_and_consume(key, window, limit)
```

Just with a **cluster client** instead of a single-node client.

### Option B – Separate Redis for rate limiting

Sometimes teams dedicate a separate Redis cluster just for rate limiting, so:

- Rate limiter is isolated from app/session/cache workloads.
- Tuning (maxmemory, eviction policy, etc.) is specific to rate data.

### 5. Tweaks if we anticipate huge scale

If we expect **tens of millions** of active users or more, we can further reduce pressure:

1. **Use coarser buckets** instead of per-request logs
  - Store counts per minute in a HASH or simple key.
  - Approximate sliding window instead of exact log.
2. **Smaller windows or lower limits**
  - Fewer timestamps per key.
3. **Per-tenant sampling or hierarchical limits**
  - Strong limits per tenant; looser per user inside tenant.